

嵌入式微處理機系統

04_CPU_s

國立台灣科技大學電機系

王乃堅老師


Feb. 2024

CPU_s

- Input and output
- Supervisor mode, exceptions, traps
- Co-processors
- Memory management and address translation
- Cache
- How architecture affects program performance
- How architecture affects program power consumption

I/O devices

- Usually includes some non-digital component.
- Typical digital interface to CPU:




- **Data registers:** hold values that are treated as data by the device, such as the data read or written by a disk
- **Status registers:** provide information about the device's operation, such as whether the current transaction has completed

Application: 8251 UART

- **Universal Asynchronous Receiver Transmitter (UART):** provides serial communication.
- 8251 functions are integrated into standard PC interface chip.
- Allows many communication parameters to be programmed.

Serial communication


- Characters are transmitted separately:



Serial communication parameters

- Baud (bit) rate (bits/sec)
- Number of bits per character (5~8)
- Parity/no parity (even or odd)
- Even/odd parity
- Length of stop bit (1, 1.5, 2 bits)

8251 CPU interface



- **Transmitter Ready**: it is ready to accept a data character
- **Transmitter Empty**: when the UART has no character to send
- **Receiver Ready**: when the UART has a character to be read by CPU

Programming I/O

- Two types of instructions can support I/O:
 - special-purpose I/O instructions; (I/O mapped I/O)
 - memory-mapped load/store instructions.
- Intel x86 provides in, out instructions.
- Most other CPUs use memory-mapped I/O.
- I/O instructions do not preclude memory-mapped I/O.

ARM memory-mapped I/O

- Define location for device:

DEV1 EQU 0x1000

- Read/write code:

LDR r1, #DEV1 ; set up device address

LDR r0, [r1] ; read DEV1

LDR r0, #8 ; set up value to write

STR r0, [r1] ; write value to device

SHARC memory mapped I/O

- Device must be in external memory space (above 0x400000).

- Use DM to control access:

IO = 0x400000;

M0 = 0;

R1 = DM(IO,M0);

Peek and poke

- Traditional **HLL** (high-level language) interfaces:
- The traditional names for functions that read and write arbitrary memory locations are **peek** and **poke**.
- Peek function in C:

```
int peek(char *location) {  
    return *location; }      /* de-reference location pointer */
```

 - To read a device register

```
#define DEV1 0x1000  
  
...  
dev_status = peek(DEV1)      /* read device register */
```
- Poke function in C:

```
void poke(char *location, char newval) {  
    (*location) = newval; }   /* write to location */
```

 - To write to status register

```
poke(DEV1,8)                  /* write 8 to device register */
```



Busy-wait input and output

- Devices are typically slower than CPU
- Polling:
 - asking an I/O device whether it is whether it is finished by reading its status register
 - If CPU is writing several characters to an output device, then it must wait for one operation to complete before starting the next one.
 - If we try to start writing the second character before the device has finished with the first one, for example, the device will probably never print the first character.
 - Simplest way to program device.
 - Use instructions to test when device is ready.



Busy-wait I/O

● Busy-wait routine in C (Ex. 3-3)

```
current_char = mystring;          /* point to head of string */
while (*current_char != '\0') {    /* until null character */
    poke(OUT_CHAR,*current_char); /* send character to device */
    while (peek(OUT_STATUS) != 0); /* keep checking status */
    current_char++; /* update character pointer */
}
```

- Outer while loop sends the characters one at a time
- Inner while loop checks the device status
- By repeatedly checking the device status until the status changes to 0



Simultaneous busy/wait input and output

● Ex.3-4 Copying characters from input to output using busy-wait I/O

```
while (TRUE) {                                /* perform operation forever */
    /* read */
    while (peek(IN_STATUS) == 0); /* wait until ready */
    achar = (char)peek(IN_DATA); /* read character */
    /* write */
    poke(OUT_DATA,achar);
    poke(OUT_STATUS,1); /* turn on device */
    while (peek(OUT_STATUS) != 0); /* wait until done */
}
```

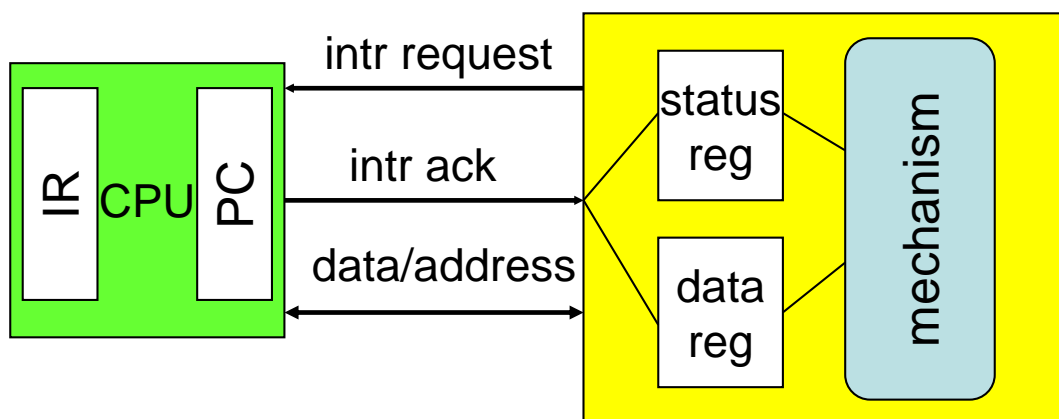


Interrupt I/O

- Busy/wait is very inefficient. (Polling)
 - CPU can't do other work while testing device
 - Hard to do simultaneous I/O → parallel processing
- CPU can do useful work in parallel with I/O transaction, such as computation, and control of other I/O
- Interrupts allow a device to change the flow of control in the CPU.
 - Causes subroutine call to handle device.



Interrupt interface



- PC → interrupt handler routine (device driver)
- Interrupt request: when it wants service from CPU
- Interrupt acknowledgement: when it is ready to handle the I/O device's request



Interrupt behavior

- Based on subroutine call mechanism.
- Interrupt forces next instruction to be a subroutine call to a predetermined location.
 - Return address is saved to resume executing **foreground program**.
- Interrupts allow the flow of control in the CPU to change easily between different contexts, such as a foreground computation and multiple I/O devices

Interrupt physical interface

- CPU and device are connected by CPU bus.
- CPU and device **handshake**:
 - device asserts **interrupt request**;
 - CPU asserts **interrupt acknowledge** when it can handle the interrupt.

Ex. 3-5: character I/O handlers


● Coping characters from input to output with basic interrupt

```
void input_handler() {           /* get a char. and put it in globle */
    achar = peek(IN_DATA);       /* get character */
    gotchar = TRUE;              /* signal to main program */
    poke(IN_STATUS,0);           /* reset status to initiate next transfer */
}
void output_handler() {          /* react to character being sent */
}
```

Ex. 3-5: interrupt-driven main program

```
main() {
    while (TRUE) {               /* read then write forever */
        if (gotchar) {           /* write a character */
            poke(OUT_DATA,achar); /* put character in device */
            poke(OUT_STATUS,1);   /* set status to initiate write */
            gotchar = FALSE;      /* reset flag */
        }
    }
}
```

● Queue for characters:




Buffer-based input handler

```
void input_handler() {  
    char achar;  
    if (full_buffer()) error = 1;  
    else { achar = peek(IN_DATA);  
          add_char(achar); }  
    poke(IN_STATUS,0);  
    if (nchars == 1)  
        { poke(OUT_DATA,remove_char());  
          poke(OUT_STATUS,1); }  
}
```



I/O sequence diagram



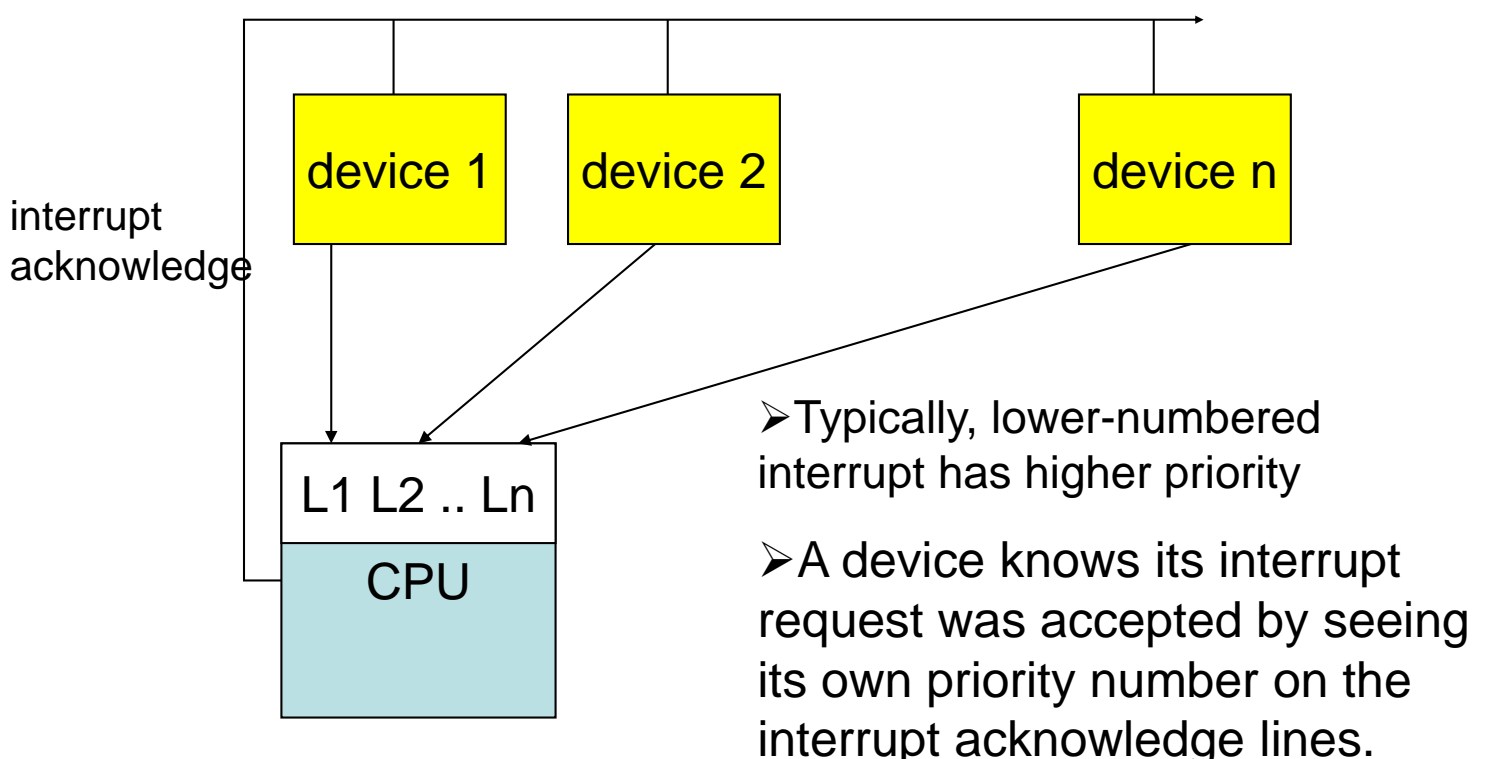
Debugging interrupt code

- What if you forget to change registers?
 - Foreground program can exhibit mysterious bugs.
 - Bugs will be hard to repeat---depend on interrupt timing.

Priorities and vectors

- Two mechanisms allow us to make interrupts more specific:
 - **Priorities** determine what interrupt gets CPU first.
 - **Vectors** determine what code is called for each type of interrupt.
- Typically, **lower**-numbered interrupt has **higher** priority
- A device knows its interrupt request was accepted by seeing its own priority number on the interrupt acknowledge lines.
- Mechanisms are orthogonal: most CPUs provide both.

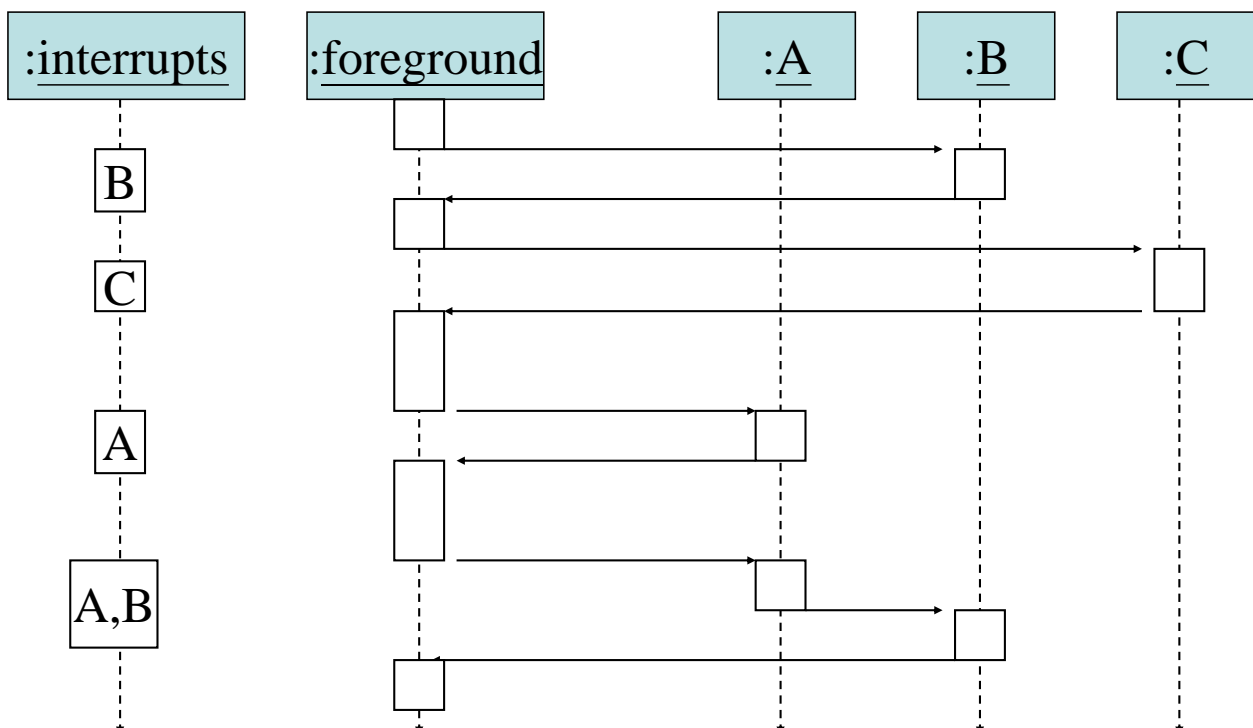
Prioritized interrupts



Interrupt prioritization


- **Masking**: interrupt with priority lower than current priority is not recognized until pending interrupt is complete.
- **Non-maskable interrupt (NMI)**: highest-priority, never masked.
 - Often used for power-down.
- We can combine polling with prioritized interrupts to efficiently handle the devices. (see Fig. 3-4)

Example 3-8: Prioritized I/O

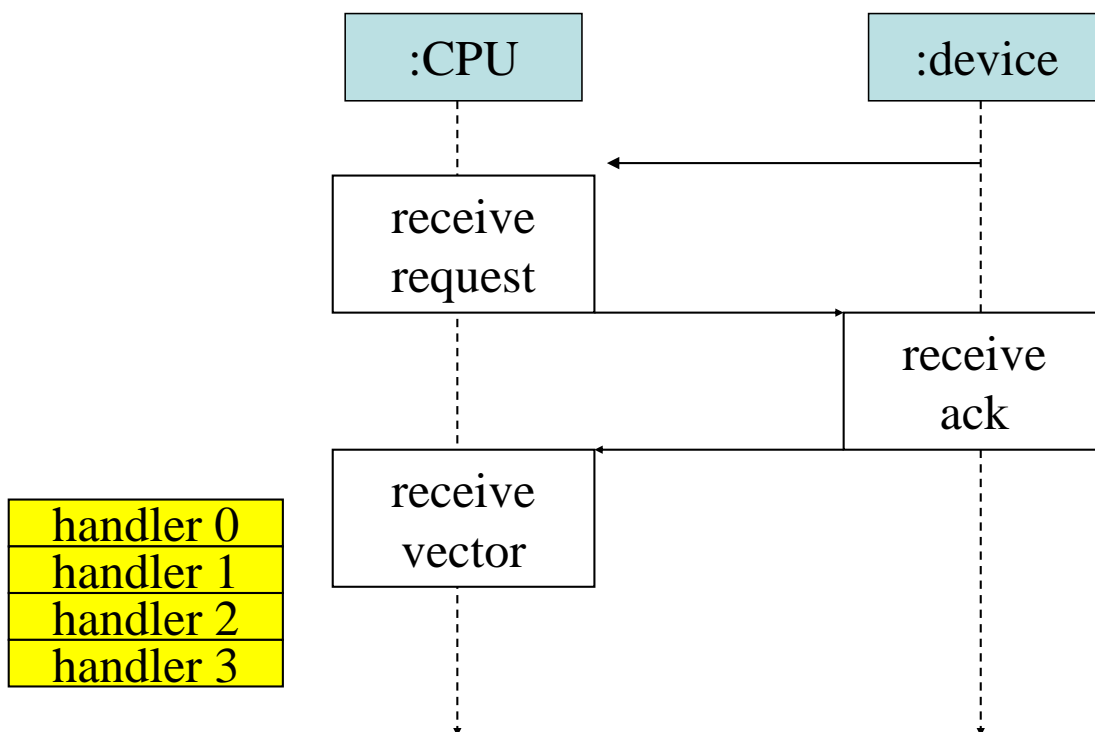


Interrupt vectors

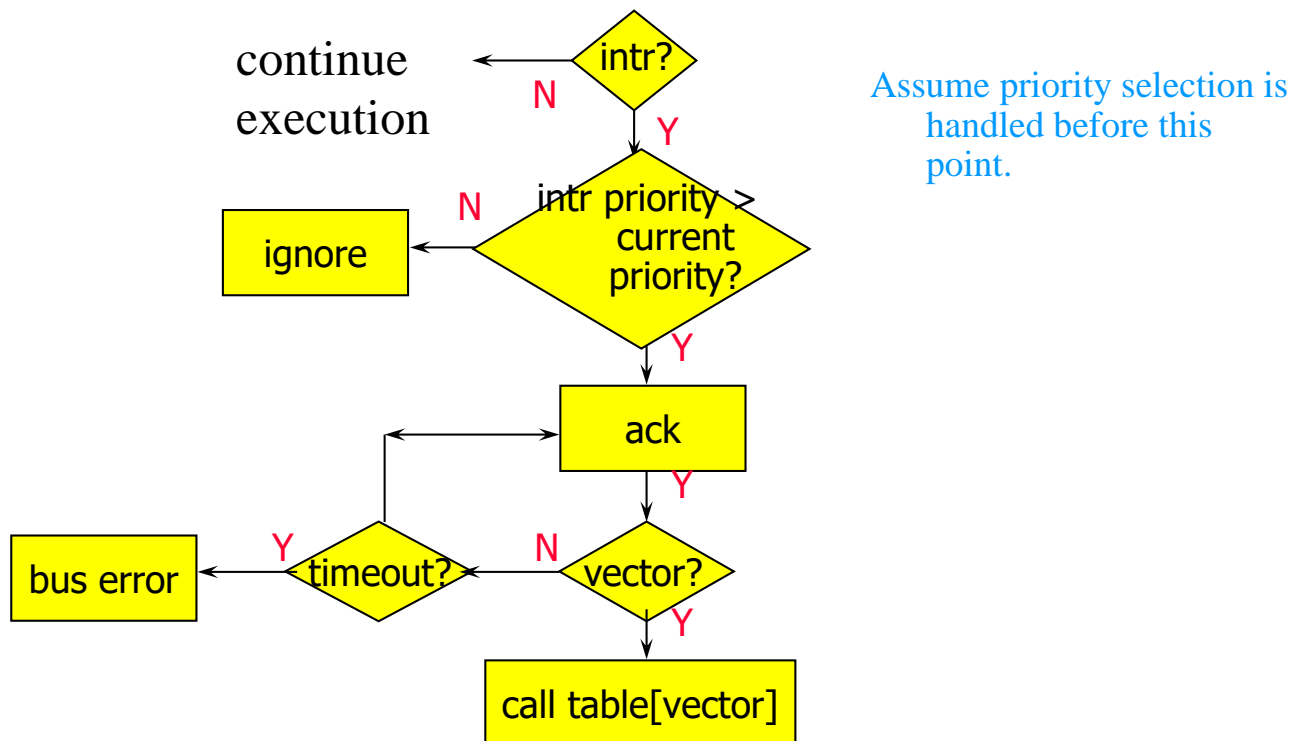
- Allow different devices to be handled by different code.
- Interrupt vector table:



Interrupt vector acquisition



Generic interrupt mechanism



Interrupt sequence

1. CPU: check pending interrupts at the beginning of an instruction. It answer the highest priority interrupt by acknowledging request.
2. Device: receives acknowledgment and sends vector.
3. CPU calls handler by LUT. A subroutine-like mechanism is used to save the PC and CPU's state.
4. Software: the device driver may save additional CPU state
 - Processes request
 - Restore the saved state
 - Interrupt return
5. CPU restores PC and state to foreground program.



Sources of interrupt overhead

- Handler execution time (similar to a subroutine)
 - Branch penalty
 - Store CPU register
- Interrupt mechanism overhead
 - Acknowledge the interrupt
 - Obtain the vector from the device
- Register save/restore.
- Pipeline-related penalties.
- Cache-related penalties.



ARM interrupts

- ARM7 supports two types of interrupts:
 - Fast interrupt requests (FIQs).
 - Interrupt requests (IRQs).
- Interrupt table starts at location 0.



ARM interrupt procedure

● CPU actions:

- Save PC. Copy CPSR to SPSR.
- Force bits in CPSR to record interrupt.
- Force PC to vector.

● Handler responsibilities:

- Restore proper PC.
- Restore CPSR from SPSR.
- Clear interrupt disable flags.

ARM interrupt latency

● Worst-case latency to respond to interrupt is 27 cycles:

- Two cycles to synchronize external request.
- Up to 20 cycles to complete current instruction.
- Three cycles for data abort.
- Two cycles to enter interrupt handling state.

Supervisor mode

- May want to provide protective barriers between programs.
 - Avoid memory corruption.
- Need **supervisor mode** to manage the various programs.
- SHARC does not have a supervisor mode.



ARM supervisor mode

- Use SWI instruction to enter supervisor mode, similar to subroutine:
SWI CODE_1
- Sets PC to 0x08.
- The bottom 5 bits of the CPSR are all set to 1
- Argument to SWI is passed to supervisor mode code.
- Saves CPSR in SPSR.



Exception

- **Exception**: internally detected error.
- Exceptions are synchronous with instructions but unpredictable.
- Build exception mechanism on top of interrupt mechanism.
- Exceptions are usually prioritized and vectorized.



Trap

- **Trap (software interrupt)**: an exception generated by an instruction.
 - Call supervisor mode.
- ARM uses SWI instruction for traps.
- SHARC offers three levels of software interrupts.
 - Called by setting bits in IRPTL register.



Co-processor

- **Co-processor**: added function unit that is called by instruction.
 - Floating-point units are often structured as co-processors.
- ARM allows up to 16 designer-selected co-processors.
 - Floating-point co-processor uses units 1 and 2.

